On dynamic extensions of context-dependent parser
— Extended Abstract —

Jūlija Pečerska[1], Artjoms Šinkarovs[2], Pavels Zaičenkovs[3]

[1] University of Latvia, Raiņa bulvāris 19, Rīga, Latvija, LV-1586
[2] Heriot-Watt University, Riccarton, Edinburgh, EH14 4AS, United Kingdom
[3] Moscow Institute of Physics and Technology, 141700, 9, Institutskii per., Dolgoprudny, Moscow Region, Russia

Abstract. The idea of using a preprocessor before the actual compilation is proven to be useful, however, the automatic verification of such a programs is extremely hard. The main difficulty comes from conceptual separation of grammatical rules of the language and substitution mechanisms underneath a preprocessor. C/C++ preprocessor is included into the standard of the languages, however, expressiveness of this tool is very limited in both computational power, and syntax extension capabilities. The modern programming languages nowadays define a syntax which cannot be described precisely using context free grammars. This problem is well known, and as a solution there are parser generators like ANTLR which generates LL(*) parsers rather than LALR/LR. Nevertheless, a lot of real-world parsers are implemented by hand.

In this paper we present a way of building a preprocessor which allows dynamic changes of the grammar and which works on top of any recursive descent parser which meets a certain requirements. The preprocessing step consists of matching a list of tokens against a regular expression built on tokens and encoded production names of the given grammar and transforming the result using a functional language. We demonstrate that using the proposed approach it becomes possible to i) give a static guarantees regarding the preprocessing rules; ii) safely express non-trivial syntactical constructions; and iii) perform a restricted partial evaluation.

1 Introduction

Very often expressiveness of programming language introduce a number ambiguities in its syntax. The language specification clearly states how to resolve the conflict, however it may not be possible to formulate the resolution in terms of context free grammars. In order to illustrate that consider the following examples.

1. The classical example from C language is a type-cast syntax. As a user can define an arbitrary type using `typedef` construct, the type casting expression `(x) + 5` is undecidable, unless we know if `x` is a type or not.

2. Assume that we extend C syntax to allow an array concatenation using infix binary `++` operator and constant-arrays to be written as `[1, 2, 3]`. We immediately run into the problem to disambiguate the following expression: `a ++ [1]`, as it could mean an application of postfix `++` indexed by `1` or it could be an array concatenation of `a` and `[1]`.
3. Assuming the language allows any unary function or operator to be applied as infix and postfix, we cannot disambiguate the following expression:

```
log (x) - log (y)
```

Potential interpretations are: `log (uminus (log (x))) (y)`, which is obviously an error, or `minus (log (x), log (y))`.

Sometimes the context may influences not only parsing decisions but also lexing decisions. Consider the following examples:

1. C++ allows nested templates, which means that one could write an expression `template <typename foo, list <int>>`, assuming that the last `>>` is two closing groups. In order to do that, the lexer must be aware of this context, as in a standard context character sequence `>>` means shift right operation.
2. Assuming that a programmer is allowed to define her own operators, the lexer rules must be changed, in case the name of a new operator extends the existing one. For example, assume one defines an operation `+-`. It means that from now on an expression `+-5` should be lexed as (`+-`, `5`), rather than (`+`, `-`, `5`).

In order to resolve the above ambiguities using table-based parser generator, we have to make sure that one can annotate the grammar with a correct choices for each shift/reduce or reduce/reduce conflict, which puts a number of requirements on the syntax of a parser generator and on the finite-state machine execution engine. Secondly, we have to introduce contexts without interfering with the above conflict-resolution. Finally, one has to have an interface to a lexer, in case lexing is context-dependent, and all the mechanisms should be aware of error-recovery facilities.

Having said that, we may see that using parser generators could be of the same challenge as writing a parser by hands, where all the ambiguities could be carefully resolved according to the language specification. As it turns out, most of the real-world language front-ends use hand-written recursive descendent parsers, specially treating ambiguous cases. For example the following languages do: C/C++/ObjectiveC in GNU GCC [2], clang in LLVM [4], JavaScript in Google V8 [3].

The main goal of any preprocessor is to perform a substitution of one sequence with another. The unit of the sequence may be different depending on the agreement, however the common case is to say that the unit is a sequence of characters of the same class. The number of classes is normally fixed, however character belonging to the class may be static as in C preprocessor, where

for example notion of space cannot be changed; or dynamic as in TEX, where one could specify that a certain character is a delimiter. Then the substitution itself is a replacement of units in a sequence which are treated as arguments, with the assigned arguments. The key problem here, as we are concerned in this paper is a lack of separation between the rewriting itself and transformation of token-sequence. Consider an example of a C macro:

```
#define foo(x, y) x y
```

First of all, it is really hard to say anything about the result of this macro, as `foo (5,6)` expands to `5 6`, but both `foo (,5)` and `foo (5,)` expands to `5`. Secondly, as comma is a part of syntax definition of the macro then one cannot just pass a token sequence `5, 6` as a first argument of `foo`. In order to resolve this one may escape the comma by wrapping it in parentheses and calling `foo ((5,6), 7)` which will expand to `(5, 6) 7`. The only way to flattern the list is to perform an application of another macro. For example:

```
#define first(x, y) x
#define bar(x, y) first x y
```

So we have higher-order macro here, but it would work only if arguments have a right type and the application of `bar ((5,6), x)` would expand to `5, x`, however application of `bar (5, 6)` would not provide an error but would expand to `first 5 6`. And as a last example we can make original macro `foo` return 3 arguments by expanding `foo (5, foo (6, 7))` which will expand to `5 6 7`. We may clearly see that making some static conclusions by checking a system of macros is impossible, as it may all depend on the application; and making any dynamic decisions with respect to the correctness of substitution is also not possible, as there is no way to declare the criteria of correctness.

Despite all the correctness complications macro systems are are not powerful enough to introduce a new language constructs. For example, it would be natural to write an absolute value as `|a|`, or to allow a number of user-defined literals to introduce units in a programming language like `5 kg` or `8 mm`. Even if a macro-system can do it, then it is may be confusing to resolve nested expressions treating one and the same symbol differently. For example, would it be possible for some macro system to transform an expression `| a|b |` into `abs (a|b)`?

The proper way of doing macro-substitutions is to allow an extension to the grammar. However, providing a handle on arbitrary changes of the grammar may lead to uncontrolled changes in the semantics of the language, which again would be hard to prove if the program is correct.

As a solution to the given problem we introduce a preprocessing which works on the sequence of pseudo-tokens which is a combination of tokens recognized by parser and productions of the grammar used by parser. To make it even more powerful we allow a regular expression on pseudo-tokens, still being able to guarantee a correctness of the transformation.

The rest of the paper is organized as following:

FIXME: *BLA-bla-bla*

2 Parser model

Our work is concerned with a dynamic grammar modification on the fly, and as a base of our approach we are going to consider an LL(k) recursive descent parser with a certain properties.

As a running example in this paper we are going to use an imaginary language with a C-like syntax. Consider a grammar of the language.

```
program       ::=  ( function ) * ;
function      ::=  type-id  '(' arg-list ')' stmt-block ;
arg-list      ::=  ( type-id  id ) * ;
stmt-block    ::=  '{' ( expr | return ';' ) * '}' ;
expr          ::=  fun-call | assign | cond-expr ;
fun-call      ::=  id '(' ( expr ) * ')' ;
assign        ::=  id '=' expr ;
cond-expr     ::=  bin-expr '?' cond-expr ':' expr ;
bin-expr      ::=  bin-expr binop primary-expr
primary-expr  ::=  number | prefix-op expr | '(' expr ')' ;
binop         ::=  '&&' | '||' | '==' | '!=' ... ;
prefix-op     ::=  '-' | '+' | '!' | '~' ;
```

First of all we ask, that every production is represented as a function with a signature `Parser -> (AST|Error)`, i.e. function gets a parser-object on input and returns either an AST node or an error. We would call those functions handle-functions. We require that handle-functions structure mimic a formulation of the grammar, i.e. if a production A depends on a production B, we require function handle-A to call function handle-B.

Each handle-function implements error recovery (if needed) and takes care about disambiguating productions according to the language specification, resolving operation priorities, syntax ambiguities and so on. Each handle function has an access to the parser, which keeps has an internal state, which changes when a handle-function is applied. In a some sense an application of a handle-function is a reduce step of a shift-reducer.

Each handle-function is paired with a predicate function which checks whether a sequence of tokens pointed by a parser-state matches a given rule. This type of functions we will call is-functions. Application of an is-function does not modify the state of the parser. Is-functions may require unbounded look-ahead from the parser, which also happens to be a requirement. We assume that in order to resolve complicated ambiguities unbounded look-ahead is needed anyways, as language expressions normally allow unbounded nesting.

Assuming that all the requirements are met, the grammar $G = (N, T, P, S)$ provides a full information required to build a support for user-defined matches.

3 Dynamic extension

We introduce a generic syntax extension which can be applied to any language recognized by a parser which meet all the requirements from section 2. The syn-

tax extension is capable to perform standard preprocessing tasks providing also a functionality to do partial evaluation and non-trivial generic code transformations.

On the user level we introduce a single macro definition which is called `match` and which substitutes a sequence of tokens matched with a certain pattern with another sequence of tokens. Consider the following example:

```
match [\expr] foo ( a , b )   -> [\expr] a + b
```

where we substitute a sequence of tokens `foo (a,b)`, which would be normally matched by an 'expr' rule of our grammar, with token-sequence `a + b` and applying 'expr' production on them. The above definition has a number of differences from the classical C preprocessor macro-definition `#define foo(a, b) a + b`:

- The above macro definition is not a function and `a` `b` are not arguments. The macro will match expressions where identifiers `a` and `b` are passed. In terms of tokens, only the sequence of tokens 'foo', '(', 'a', 'b', ')' will be matched. Hence, the match would not replace expressions `foo (2, 3)` or `foo (b, a)`.
- The match is bounded to one particular production in the grammar, which is 'expr' in this example. It means that it would not perform a substitution in case one wrote `foo (a, b)` as a member of a statement block or a function header.
- The result of the substitution is always a single value, which avoids the classical situation with missing parentheses in the macro definition, i.e. if a macro-definition `#define foo(a, b) a + b` is applied to `foo (2,3) * 5`, expansion would make it `2 + 3 * 5`, where a conceptual expansion of the above match would look like `(2+3) * 5`.

It should be pointed out that in order to associate macro with some production it's necessary to provide grammar rules for a programmer. This will allow to take context into account and to interact with grammar parser dynamically.

3.1 Language patterns

FIXME: *This subsection has to be rewritten!* The depicted parser would be impractical without pattern matching. To illustrate this we would like to match expression `foo(a, b)`, which can be occurred in place of 'expr' production and `a, b` are allowed to be any relevant arguments.

```
match [\expr] foo ( \expr , \expr )
   -> [\expr] \expr[1] + \expr[2]
```

Let's compare this example with the previous one. We state here that we expect two token sequences in the brackets that would be interpreted as 'expr' productions. The type of production is important as this allows to perform an effective type checking. Specifically, this macro will recognize `foo (return 0, 1)` as a

fallacious, unlike the C macro which will not point out any error.

It is noteworthy to mention about pitfalls of this approach. The macro extension associates user-defined rules with the grammar of the language. Therefore, these rules might conflict with existing ones and an ambiguous grammar can be produced. We state here that user has to control such situations himself, otherwise, an error of the parser will be raise.

Furthermore, we provide an interface to a lexer. For instance, it's possible to use some specific tokens in user-defined productions.

```
match [\expr] | \expr | -> [\expr] absolute_value (\expr[1])
```

Here we introduce a new | token which could be used for getting an absolute number value. A remarkable point is that expressions such as |-5| or even ||-5||, as this macro takes lexical scope into account. Notice that |1|2| will produce an error as expected.

New tokens defined in the left part of the matcher are appended to a valid token table. We can use them equally well as 'native grammar' tokens. It allows to build legacy rules using new tokens:

```
match [\expr] < \expr > -> [\expr] | \expr[1] |
```

As a matter of fact only defined tokens can be used in the right part of the macro.

FIXME: *Rewritten until this very moment*

Matches work as a standard Term Rewrite System (S, R), where S is a set of terms; $S = L_t \cup P_t$, where L_t is a set of tokens recognized by lexer, and P_t is a set of pseudo-tokens which are escaped production-names of the parser. R is a set of rewrite rules, where $\forall r_i \in R \Rightarrow r_i :: \{S \cup R_t\}^n \to S^m$. R_t is a set of regular expression symbols which is allowed to formulate a rule.

Each rule has a general form of:

$$s_1 \ \text{if} \ p(s_1) \to s_2.$$

Where s_1 is a regular expression over tokens and pseudo-tokens, $p(s_1)$ is a predicate which must evaluate to true, in order to enable match expand; s_2 is a sequence of tokens and pseudo-tokens which is used as a substitution.

The left hand side of each match allows a user to build a new production of the grammar restricted by a power of regular expressions. So, for instance it would not be possible to build a rule $(a^k b^m, k = m)$, as the new rule is recognized by DFA without any memory.

Now as a left-hand side of the match has a free form and the rewriting system is recursive by its nature, we face face a number of problems in case we want to prove correctness of the system. We have a standard word problem and stopping problem of the rewrite system. It is important to understand that in our case we are not dealing with a single TRS, but we have a mechanism to construct an arbitrary rewrite system.

Another important problem is to guarantee that the right-hand side of the match, is a valid rule in a given production of a given grammar. For example:

```
...
match [\expr] bar ( \id, \expr )      -> [\expr] \id{1} ( \expr{1} )
match [\expr] foo ( \expr , \expr )  -> [\expr] bar ( baz , \expr{1} )
match [\expr] \id ( \expr )          -> [\expr] \id{1} ( \expr{1}, \expr{1} )
...
```

Here we can see, that there is no chance to check statically whether the last rule is going to be expanded to the one of the matches or not, as it depends only on the value of `id`. Now, in case of `foo` we cannot check, if a user meant to pass a higher-order function `baz` or just a token `baz`.

It would be possible to resolve the situation at runtime, and one still would not be able to get a program that does not belong to the language generated by the grammar. However, the meaning of the program is unprovable, which practically means we have a powerful tool to obfuscate the code.

In order to resolve the situation, we want to introduce more static knowledge to the rewriting rules by introducing types. Now, we have to understand, that there is a clear distinction between the match that performs a substitution, and the helper-matches which make a transformation of the tokens matched by the left part of the match. The main reason here is that when we express a transformation of the matched token-sequence, the intermediate results we pass through helper-matches don't have to be a valid parser expressions. However we still want to type-check them. Consider the following example: (`expr + expr ...`):

```
match [\expr] foo (\expr \(, \expr \) *)
   -> [\expr] bar replace (tail (res), \, , +)
```

This match replaces function `foo (1,2,3)` call with `bar (1+2+3)`. In order to do that we want to have a generic function that operates on the list of tokens and pseudo-tokens, which replaces every occurrence of ',' with '+'. Evaluation of 'replace' happens outside of any productions and in general case, the return type of such a type of a function could be not a valid input for any parser production.

It means that here we would like to make a clear cut between the match and the match-function. We should keep in mind, that it is always possible to express any match-function using the rewriting system of the matches, however first of all we would like to introduce the semantics and type system description for both matches and match-functions.

3.2   Type system

Obviously matches and match-functions have to share the notion of types they are operating with. What kind of types the matched left-hand side of the match can produce? If we consider that each pseudo-token represents a type, we may note that the regular expression automatically generates an algebraic data type. This is fairly easy to prove constructively:

1. Each pseudo-token generates a type.
2. Each concatenation generates a tuple.

3. Each choice generates an alternative. For instance \id|\expr can be represented with a type (id|expr).
4. Each asterisk generates a list of types generated by a sub-asterisk expression. For example: \(a | b \) * can be represented as [(a | b)].

Finally, in order to express a bounded recursion, one has to operate with integer and boolean types. Supporting algebraic data types comes with several built-in function in order to traverse the lists and to find the type of the variable at runtime, for being able to branch depending on the type of the expression. So we introduce the following built-in functions:

head Returns the first element of the list or nil
tail Returns the list without the first element.
concat Construct a list from two lists.
type Return a type of a given expression.

Finally we have to introduce the syntax of the match-functions and describe the way one can check that the type returned by an application of the match-functions is correct with respect to the grammar rule on the right hand side of the match.

3.3 Match-function syntax

The syntax of the match function can be derived from a functional language like ML, making sure that the types are properly recognized.

FIXME: *BLA-BLA-BLA*
Now let's consider an example which replaces foo (expr, expr ...) with bar (expr + expr + ...).

```
match-fun replace :: ([expr|\op[',']], \op[')']) \op \op
                   -> ([expr|\op['+']], \op[')'])
replace lst r w = lst if len (lst) == 1
                  |
                  concat (head (lst), replace (tail (lst), r, w))
                  if (type (head (lst)) == \expr)
                  |
                  tail (lst) otherwise

match [\expr] foo ( \expr \(, \expr\) * )
   -> [\expr] bar ( replace (tail (tail (res)), \,, + )
```

4 Regular expressions

FIXME: *This is a weird draft by Petch*

4.1 Match as a regular expression

The left part of the match (without the resulting type) is actually a regular expression, but with tokens to be matched instead of single characters. That does not change the concept, because just the same as we can have a getter for the next symbol in the input stream, we have the get next token function in the parser. We operate not with an input stream of symbols, but with an input stream of pseudo-tokens.

Currently our regular expression syntax allows using or | and asterisk *
notations. Asterisk is more binding than or, so if you want to have (a or b) zero to n times you will have to write `(a|b)*`. The supported syntax can easily be extended. For the time matching classes of tokens is unsupported, however the classes can be emulated by using or constructions. This is possible because the token count is fairly limited, unlike the character set that can be used in regular expressions.

The given regular expression is parsed to create a non-deterministic automaton.

4.2 NFA to DFA

Next step is to create a DFA out of the created NFA to minimize the effort needed to execute the given automata. Currently the subset construction algorithm is used for the determinisation process. It is described in detail in the "Dragon Book".

Then the determinate automata created is minimised. Minimisation algorithm is, again, described in detail in the "Dragon Book". The algorithm creates sets of states that cannot be distinguished by any input token sequence. Once the algorithm fails to break the sets into smaller ones it stops the process. These sets of states then become the new states of the minimal automaton.

It is essential to note that the existence of such minimal automaton is provable, despite the complexity of the regular expression it describes. This statement implies that we have a possibility to prove equality of two automata. As the naming of the states is unimportant, we will say that two automata are the same up to state names if one can be transformed to another be simply renaming the states. Therefore two regular expressions match the same input if and only if their automata are the same up to state names.

The created DFA is minimal and therefore the most effective for matching the given expression. The DFA is represented by a list of objects where each object contains a map of symbols and objects to which the current symbol transfers the automaton.

The construction of the DFA does not support back-referencing, so the bracket groups are intended only to change the priority of the operations in the regular expressions. We abandon the back-referencing feature in favor of maintaining a fully determinate automaton for each regular expression. Consequently, in the matches' output, all of the tokens are returned in a single unnested list.

FIXME: *is the paragraph below needed?*

We still consider an option to allow backtracking in case it will become necessary. If we perform determinisation only on the parts of the automaton that are enclosed in braces and then combine the created automata by consequently connecting the accepting and starting states together, we will create an automaton that allows grouping. However this approach disallows the combining of several matcher automata into a single one, as the procedure would ruin the grouping.

4.3 Several DFA to single DFA

As we operate with the source code, we want to check all of the matches that we have in one pass through the code. We will consider two possibilities to do this, which are described below. For both of them we evaluate DFA adding, matching and context inheritance algorithmic difficulty.

Context inheritance difficulty is important, as we can have several included contexts, where the matches introduced inside the included one have a higher priority. Imagining that we have a system of `n` automata we have two options of doing so. The first one is that on entering a context we add the `m` matches we come across to the main match set and remove the `m` matches upon exiting the context. The second one is that on entering a context we create a copy of the parent context, to which we add the `m` newly found matches and upon exiting the context throw out the entire new system of matches.

We do not explicitly select a method of operation with joining DFA, we only give the complexities of the proposed variants. This is because the optimal solution selection should be based on the practical uses of the system. It is impossible to state, without any actual examples, what will be more time-effective during execution. Even though the second option is time-consuming when adding and removing matches, the dramatic improvement in execution time might come from the fact that match count is relatively small, but the automaton execution will, at worst cases, be performed for every token in the source text.

List of DFA The simplest method of executing several DFA at one pass though the source code is to store them in a list. Let us assume that there are `n` matches to be checked. Then the automata list represents an NFA with `n` epsilon branches from the start state, each of which leads to one of the DFA we already created.

Match adding

DFA adding in this case is simple, as the only action necessary is to add the automaton to the list, so the complexity is $O(1)$.

Matching

In this case matching depends both on the count of automata `n` and on the length of the matched pattern `l`, as we have `n` branches of an NFA to traverse simultaneously. So in the worst case the algorithmic difficulty of the task will be $O(n * l)$.

Context inheritance

The first option of context inheritance can be executed easily, as the added `m` matches can be removed from the system in a single block. So it's complexity is $O(1)$. The creation and deletion of the whole match system also has time complexity of $O(1)$.

Combining several DFA into single DFA In this case we combine the automata created for each of the matches together into a single DFA. This is done in order to reduce matching time.

Match adding

The algorithm combines the states of two automata one by one, and is, in fact, an adaptation of the determinisation algorithm. Its complexity is $O(m*n)$, where `m` is the amount of states in one automata, and `n` is the amount of states in the second automata.

Imagine that we have two consequent matches, for both of which we have generated an automaton. At first we create a starting state for the merged automaton by combining the starting states of existing ones. Then we examine the paths leading by same tokens from the combined states in the two starting automata. We create states in the new automata from the end points of the paths in the existing automata and add paths to them by the examined symbol. By doing so for every new state we add all of the paths and create the combined automaton. By prioritising the regular expressions we can unambiguously identify the regular expression for which a specific state will be accepting.

The figures show an example of this algorithm execution. Figure 1 shows the automaton for `a*b*`. It has two states, each of which is an accepting state. Figure 2, on the other hand, shows the automaton for regular expression `ab`. It has 3 states, the last of which is an accepting state. Smaller circles in colors of the state id's show that the state is accepting for the corresponding automata.


Fig. 1. Automaton for a*b*

The algorithm traverses the two automata step by step adding new states to the resulting one. Firstly, the state `(1, 1)` is created, which is a combination

Fig. 2. Automaton for ab

of the two automata starting states. By symbol `a` the first automata goes to state `2`, however the second one stays at state `1`. As a result state `(2, 1)` is added together with a path to it by token `a` from `(1, 1)`. Next paths are created for the newly added state `(2, 1)`. By token `b` we add state `(3, 2)` and the corresponding path and so on. The resulting automata is shown on figure 3.

Fig. 3. Automaton for ab and a*b*

In figure 3 the states `(1, 1)`, `(2, 1)`, `(1)` and `(2)` are accepting for the regular expression `a*b*`, but the state `(3, 2)` is accepting for regular expression `ab`. Here we assume that the first regular expression to appear has lower priority than the ones that appear later. Otherwise state `(3, 2)` would also be accepting, but for the second automata.

Matching

Here matching depends only on the length of the matched token list, so the complexity is $O(l)$, as only a single automaton traverse option exists for a single input token.

Context inheritance

Context inheritance by adding and removing matches is not as efficient as in the case of automata list. Adding complexity was described above, however

to remove an amtch from the system we will have to traverse all of the automata states and track those, that belong to the removed automata. So the complexity of a removing of an expression, if `m` is the amount of states of the system, will be $O(m)$.

If we adopt the second option, creation and deletion of a match system for a nested context, the complexity will remain the same as in adding DFA to the system, as the deletion will not add to it.

5 Application

5.1 Preprocessing

5.2 Templates

5.3 Optimisation potential

6 Evaluation

Here is a bunch of links for the existing macro-preprocessors:

ML/I	http://www.ml1.org.uk/htmldoc/ml1sig.html
GEMA	http://gema.sourceforge.net/new/docs.shtml
GPP	http://files.nothingisreal.com/software/gpp/gpp.html В этой штуке советую заглянуть в ADVANCED EXAMPLES с лямбдой
TRAC	http://web.archive.org/web/20050205172849/http://tracfoundation.org/t2001tech.htm Это очень разумная идея правда совсем дохлая – там тоже функциональный язык внутри живет, но работает на строках кажись

Еще бывают: m4, cpp, lisp/scheme macros, tex?...

6.1 Macros in Lisp

Due to the Lisp's fully parenthesized Polish prefix syntax notation there is a powerful macro engine. We are going to cover its main principles and advantages.

1. There is no need to mark out a structure from a token sequence for the internal representation (IR). Every sequence of tokens in parenthesis shapes an expression called *form*. A program on Lisp represents a tree of nested forms. This property is called *homoiconicity*, as the source code of a program can be proceeded by means of the same language. In this case it is possible to consider a macro-definition as a left tree substitution for one in the right part. For example,

```
(defmacro sum (x y z)
(list '+ (list '* x y) (list '* z z))
```

macros defines a list substitution for the nested list structure.

2. Any valid expression in Lisp represents a form. Even the whole program is a form. Therefore, it is possible to replace huge parts of a program and even the whole program.

3. Lisp's compiler does not match the whole expression (or a form) in order to find out either a corresponding macros is defined. Due to the prefix notation a 'keyword' is the first token in a form. If a macro definition is occured in the macro table under the same keyword, then we have to perform a macro substitution.

   However, this simplicity has a disadvantage too. Macros overloading is not allowed in Lisp.

4. A macro processor in Lisp not only substitutes expressions, but also evaluates them. There is an expression-value table in Lisp, where each expression is associated with it's value. For example, if we write `(setq a 3)`, then an expression `a` has a corresponding value `3`. Consequently, if we create a list `(list a (+ a 1))`, a list `(3 4)` will be returned, as this expression will be evaluated during compilation. This emphasizes an interpetive nature of the language. Although expression evaluation is not covered in the Common Lisp standart, many compilers, such as GNU CLisp, CMU Common List support this feature.

7   ML/1 macros

ML/1 is a stream-based macros processor[1]. It operates on a sequence of tokens. The processor reads tokens one by one and performs input stream transformation taking into account the rules defined.

In macro definition we define a correspondence between token sequence from the input and replace tokens. All atomic tokens are separated by delimiter tokens, such as 'space' or 'new line'. Suprisingly, there are no arguments placed in the macro rule to match. The arguments can appear between any atomic tokens. For example,

```
MCDEF foo bar baz AS ...
```

will match `foo xx bar yy baz`, because arguments are inserted between atomic tokens. In order to restrict such insertion, tokens have to be combined into another one, atomic token. Thus in most cases the information about exact number of arguments and their names is not accessible. Therefore, arguments are accessed by number in the order they are met in the input string. Basically, this allows to support variable number of arguments. Here arises a problem of handling these arguments, because we do not know in advance how many arguments we have and we don't even know their types.

The following features of macro language allow to handle arguments properly:

1. Tokens placed between argument tokens are called delimiters. It is possible not only to access arguments, but also delimiter tokens, enumerating them.

2. It is possible to define 'if' condition statements. 'Jumps' or 'goto' statemens are also supported. Consequently, we can verify delimiter tokens number and its type, and perform substitutions accordingly.
3. Local variables can be used inside macros. This allows to describe loop statements for iterating over the arguments.
4. ML/1 supports nested macro calls. While searching for delimiters and arguments we can meet another macro call. In this case, we descend to a lower level and return it's delimiters and arguments. Finally, they are inserted in the 'top' list.

To sum up, ML/1 provides advanced features for macro processing. It is implemented as an imperative language operating on the stream of tokens. It supports conditions, loops, branchs and assignment, so the language is Turing-complete.

8   Future work

References

1. Bob Eager. The ML/I macro processor. http://www.ml1.org.uk/.
2. Free Software Foundation. gcc. http://gcc.gnu.org.
3. Google. V8 JavaScript Engine. htthttp://code.google.com/p/v8/.
4. LLVM Team. clang. http://clang.llvm.org.